

Residential Electricity Consumption Forecasting (CNN–BiLSTM–Self Attention)

Dataset: Individual Household Electric Power Consumption (UCI, id=235 via `ucimlrepo`)

Goal: Predict next-step `Global_active_power` using historical multivariate time-series windows.

Roadmap

1. Install + imports
2. Load dataset
3. Build clean datetime index (Date + Time)
4. Missing values
5. Resample (hourly)
6. Feature selection (MIC)
7. Scaling + windowing
8. CNN–BiLSTM–SelfAttention training
9. Evaluate + save artifacts

```
!pip -q install ucimlrepo joblib
import os, json, math, warnings
warnings.filterwarnings("ignore")

import numpy as np
import pandas as pd

import matplotlib.pyplot as plt

from sklearn.model_selection import train_test_split
from sklearn.preprocessing import MinMaxScaler
from sklearn.metrics import mean_absolute_error, mean_squared_error,
r2_score

import joblib

# TensorFlow / Keras
import tensorflow as tf
from tensorflow import keras
from tensorflow.keras import layers

print("Imports done")
print("TensorFlow:", tf.__version__)
print("GPU Available:", len(tf.config.list_physical_devices('GPU')) > 0)
```

```
2026-02-02 15:05:23.807503: E
external/local_xla/xla/stream_executor/cuda/cuda_fft.cc:467] Unable to
register cuFFT factory: Attempting to register factory for plugin
cuFFT when one has already been registered
WARNING: All log messages before absl::InitializeLog() is called are
written to STDERR
E0000 00:00:1770044724.328056      55 cuda_dnn.cc:8579] Unable to
register cuDNN factory: Attempting to register factory for plugin
cuDNN when one has already been registered
E0000 00:00:1770044724.477633      55 cuda_blas.cc:1407] Unable to
register cuBLAS factory: Attempting to register factory for plugin
cuBLAS when one has already been registered
W0000 00:00:1770044725.653282      55 computation_placer.cc:177]
computation placer already registered. Please check linkage and avoid
linking the same target more than once.
W0000 00:00:1770044725.653350      55 computation_placer.cc:177]
computation placer already registered. Please check linkage and avoid
linking the same target more than once.
W0000 00:00:1770044725.653353      55 computation_placer.cc:177]
computation placer already registered. Please check linkage and avoid
linking the same target more than once.
W0000 00:00:1770044725.653355      55 computation_placer.cc:177]
computation placer already registered. Please check linkage and avoid
linking the same target more than once.
```

□ Imports done

```
TensorFlow: 2.19.0
GPU Available: True
```

```
from ucimlrepo import fetch_ucirepo

# fetch dataset
dataset = fetch_ucirepo(id=235)

# UCI repo returns features & (sometimes) targets; this dataset is
# basically all features
df = dataset.data.features.copy()

print("□ Raw loaded shape:", df.shape)
display(df.head(5))
```

□ Raw loaded shape: (2075259, 9)

	Date	Time	Global_active_power	Global_reactive_power
Voltage	\			
0	16/12/2006	17:24:00	4.216	0.418
234.840				
1	16/12/2006	17:25:00	5.360	0.436
233.630				
2	16/12/2006	17:26:00	5.374	0.498

```

233.290
3 16/12/2006 17:27:00           5.388          0.502
233.740
4 16/12/2006 17:28:00           3.666          0.528
235.680

   Global_intensity Sub_metering_1 Sub_metering_2 Sub_metering_3
0            18.400        0.000       1.000        17.0
1            23.000        0.000       1.000        16.0
2            23.000        0.000       2.000        17.0
3            23.000        0.000       1.000        17.0
4            15.800        0.000       1.000        17.0

print("== METADATA ==")
print(dataset.metadata)

print("\n== VARIABLES ==")
display(dataset.variables)

== METADATA ==
{'uci_id': 235, 'name': 'Individual Household Electric Power Consumption', 'repository_url': 'https://archive.ics.uci.edu/dataset/235/individual+household+electric+power+consumption', 'data_url': 'https://archive.ics.uci.edu/static/public/235/data.csv', 'abstract': 'Measurements of electric power consumption in one household with a one-minute sampling rate over a period of almost 4 years. Different electrical quantities and some sub-metering values are available.', 'area': 'Physics and Chemistry', 'tasks': ['Regression', 'Clustering'], 'characteristics': ['Multivariate', 'Time-Series'], 'num_instances': 2075259, 'num_features': 9, 'feature_types': ['Real'], 'demographics': [], 'target_col': None, 'index_col': None, 'has_missing_values': 'no', 'missing_values_symbol': None, 'year_of_dataset_creation': 2006, 'last_updated': 'Fri Mar 08 2024', 'dataset_doi': '10.24432/C58K54', 'creators': ['Georges Hebrail', 'Alice Berard'], 'intro_paper': None, 'additional_info': {'summary': 'This archive contains 2075259 measurements gathered in a house located in Sceaux (7km of Paris, France) between December 2006 and November 2010 (47 months).\r\nNotes: \r\n1.\r\n(global_active_power*1000/60 - sub_metering_1 - sub_metering_2 - sub_metering_3) represents the active energy consumed every minute (in watt hour) in the household by electrical equipment not measured in sub-meterings 1, 2 and 3.\r\n2.The dataset contains some missing values in the measurements (nearly 1,25% of the rows). All calendar timestamps are present in the dataset but for some timestamps, the measurement values are missing: a missing value is represented by the absence of value between two consecutive semi-colon attribute separators. For instance, the dataset shows missing values on April 28, 2007.', 'purpose': None, 'funded_by': None, 'instances_represent': None, 'recommended_data_splits': None, 'sensitive_data': None,

```

```
'preprocessing_description': None, 'variable_info': '1.date: Date in
format dd/mm/yyyy\r\n2.time: time in format hh:mm:ss\r\n
n3.global_active_power: household global minute-averaged active power
(in kilowatt)\r\nn4.global_reactive_power: household global minute-
averaged reactive power (in kilowatt)\r\nn5.voltage: minute-averaged
voltage (in volt)\r\nn6.global_intensity: household global minute-
averaged current intensity (in ampere)\r\nn7.sub_metering_1: energy
sub-metering No. 1 (in watt-hour of active energy). It corresponds to
the kitchen, containing mainly a dishwasher, an oven and a microwave
(hot plates are not electric but gas powered).\r\nn8.sub_metering_2:
energy sub-metering No. 2 (in watt-hour of active energy). It
corresponds to the laundry room, containing a washing-machine, a
tumble-drier, a refrigerator and a light.\r\nn9.sub_metering_3: energy
sub-metering No. 3 (in watt-hour of active energy). It corresponds to
an electric water-heater and an air-conditioner.', 'citation': None}}
```

==== VARIABLES ====

		name	role	type	demographic	description
units \						
0		Date	Feature	Date	None	None
None						
1		Time	Feature	Categorical	None	None
None						
2	Global_active_power	Feature	Continuous		None	None
None						
3	Global_reactive_power	Feature	Continuous		None	None
None						
4	Voltage	Feature	Continuous		None	None
None						
5	Global_intensity	Feature	Continuous		None	None
None						
6	Sub_metering_1	Feature	Continuous		None	None
None						
7	Sub_metering_2	Feature	Continuous		None	None
None						
8	Sub_metering_3	Feature	Continuous		None	None
None						
	missing_values					
0		no				
1		no				
2		no				
3		no				
4		no				
5		no				
6		no				
7		no				
8		no				

```

# Combine Date + Time into a single datetime column
# Date format: dd/mm/yyyy, Time format: hh:mm:ss
df["datetime"] = pd.to_datetime(
    df["Date"].astype(str) + " " + df["Time"].astype(str),
    format="%d/%m/%Y %H:%M:%S",
    errors="coerce"
)

# Drop rows where datetime couldn't be parsed (should be extremely rare)
df = df.dropna(subset=["datetime"]).copy()

# Set datetime index and sort
df = df.set_index("datetime").sort_index()

# Drop original Date/Time columns
df = df.drop(columns=["Date", "Time"], errors="ignore")

# Convert all remaining columns to numeric (dataset may load as strings)
for col in df.columns:
    df[col] = pd.to_numeric(df[col], errors="coerce")

print("After datetime + numeric conversion:", df.shape)
print("Index range:", df.index.min(), "→", df.index.max())
display(df.head(5))

```

After datetime + numeric conversion: (2075259, 7)
Index range: 2006-12-16 17:24:00 → 2010-11-26 21:02:00

	Global_active_power	Global_reactive_power
Voltage \ datetime		
2006-12-16 17:24:00	4.216	0.418
234.84		
2006-12-16 17:25:00	5.360	0.436
233.63		
2006-12-16 17:26:00	5.374	0.498
233.29		
2006-12-16 17:27:00	5.388	0.502
233.74		
2006-12-16 17:28:00	3.666	0.528
235.68		
Global_intensity \ datetime	Sub_metering_1	Sub_metering_2
2006-12-16 17:24:00	18.4	0.0
		1.0

2006-12-16 17:25:00	23.0	0.0	1.0
2006-12-16 17:26:00	23.0	0.0	2.0
2006-12-16 17:27:00	23.0	0.0	1.0
2006-12-16 17:28:00	15.8	0.0	1.0

```
Sub_metering_3
datetime
2006-12-16 17:24:00      17.0
2006-12-16 17:25:00      16.0
2006-12-16 17:26:00      17.0
2006-12-16 17:27:00      17.0
2006-12-16 17:28:00      17.0

missing = df.isna().sum().sort_values(ascending=False)
print("□ Missing values per column:")
display(missing[missing > 0])

print("Total missing cells:", int(df.isna().sum().sum()))
print("Total rows:", len(df))

□ Missing values per column:

Global_active_power      25979
Global_reactive_power     25979
Voltage                   25979
Global_intensity          25979
Sub_metering_1             25979
Sub_metering_2             25979
Sub_metering_3             25979
dtype: int64

Total missing cells: 181853
Total rows: 2075259
```

Step 4 — Handle Missing Values

The dataset contains missing measurements (~1.25%).

Since timestamps are continuous and time-ordered, we use:

- Forward fill (`ffill`) to propagate recent valid measurements
- Backward fill (`bfill`) to handle missing values at the beginning of series

This preserves time continuity and is commonly used for sensor time-series.

```
# Missing value handling
before_missing = int(df.isna().sum().sum())
```

```

print("Missing cells BEFORE:", before_missing)

df = df.ffill().bfill()

after_missing = int(df.isna().sum().sum())
print("Missing cells AFTER:", after_missing)

# sanity check
assert after_missing == 0, "Missing values still exist after filling."
print("Missing values handled successfully.")

Missing cells BEFORE: 181853
Missing cells AFTER: 0
Missing values handled successfully.

```

Step 5 — Resample to Hourly

Minute-level data is large and noisy for deep learning. We resample to hourly averages:

- reduces size dramatically (faster training)
- keeps temporal patterns (daily cycles, weekly habits)
- best practical resolution for residential forecasting

We use `.resample('H').mean()` for all continuous signals.

```

# Hourly resampling (mean)
df_hourly = df.resample("H").mean()

print("Hourly shape:", df_hourly.shape)
print("Hourly index range:", df_hourly.index.min(), "->",
      df_hourly.index.max())
display(df_hourly.head(5))

# check missing introduced by resampling (rare)
print("Missing after resample:", int(df_hourly.isna().sum().sum()))
df_hourly = df_hourly.ffill().bfill()
assert int(df_hourly.isna().sum().sum()) == 0
print("Hourly dataset ready.")

Hourly shape: (34589, 7)
Hourly index range: 2006-12-16 17:00:00 → 2010-11-26 21:00:00

          Global_active_power  Global_reactive_power
Voltage \
datetime

2006-12-16 17:00:00           4.222889         0.229000
234.643889
2006-12-16 18:00:00           3.632200         0.080033

```

```

234.580167
2006-12-16 19:00:00      3.400233      0.085233
233.232500
2006-12-16 20:00:00      3.268567      0.075100
234.071500
2006-12-16 21:00:00      3.056467      0.076667
237.158667

          Global_intensity  Sub_metering_1  Sub_metering_2
\datetime
2006-12-16 17:00:00      18.100000      0.0      0.527778
2006-12-16 18:00:00      15.600000      0.0      6.716667
2006-12-16 19:00:00      14.503333      0.0      1.433333
2006-12-16 20:00:00      13.916667      0.0      0.000000
2006-12-16 21:00:00      13.046667      0.0      0.416667

          Sub_metering_3
\datetime
2006-12-16 17:00:00      16.861111
2006-12-16 18:00:00      16.866667
2006-12-16 19:00:00      16.683333
2006-12-16 20:00:00      16.783333
2006-12-16 21:00:00      17.216667

Missing after resample: 0
[] Hourly dataset ready.

```

Step 6 — Feature Selection (MIC)

We select input features using MIC (Maximal Information Coefficient), which captures both linear and non-linear relationships with the target.

Target:

- `Global_active_power`

We compute MIC score for each feature vs the target and keep the best ones.

Artifacts saved:

- `selected_features.json` (feature order is critical for Flask integration later)

```
!pip -q install minepy
```

```

from minepy import MINE

target_col = "Global_active_power"

# Features candidates = everything except target
feature_cols = [c for c in df_hourly.columns if c != target_col]

# Compute MIC
mic_scores = {}
mine = MINE(alpha=0.6, c=15)

y = df_hourly[target_col].values

for col in feature_cols:
    x = df_hourly[col].values
    mine.compute_score(x, y)
    mic_scores[col] = mine.mic()

mic_df = pd.DataFrame({"feature": list(mic_scores.keys()), "MIC": list(mic_scores.values())})
mic_df = mic_df.sort_values("MIC",
ascending=False).reset_index(drop=True)

print("MIC scores (high → important):")
display(mic_df)

# Choose best features (keep all > 0.1, or at least top 5)
selected = mic_df[mic_df["MIC"] >= 0.10]["feature"].tolist()
if len(selected) < 5:
    selected = mic_df.head(5)["feature"].tolist()

print("Selected features:", selected)

# Save selected feature list (order is important!)
os.makedirs("/kaggle/working/artifacts", exist_ok=True)
with open("/kaggle/working/artifacts/selected_features.json", "w") as f:
    json.dump(selected, f, indent=2)

print("Saved: /kaggle/working/artifacts/selected_features.json")

```

497.0/497.0 kB 7.9 MB/s eta
0:00:00a 0:00:01

etadata (setup.py) ... error: subprocess-exited-with-error

x python setup.py bdist_wheel did not run successfully.
└ exit code: 1
 > See above for output.

note: This error originates from a subprocess, and is likely not a problem with pip.

```
Building wheel for minepy (setup.py) ...    ERROR: Failed building
wheel for minepy
ERROR: ERROR: Failed to build installable wheels for some
pyproject.toml based projects (minepy)

-----
-----
ModuleNotFoundError                      Traceback (most recent call
last)
/tmp/ipykernel_55/1180116210.py in <cell line: 0>()
      1 get_ipython().system('pip -q install minepy')
      2
----> 3 from minepy import MINE
      4
      5 target_col = "Global_active_power"

ModuleNotFoundError: No module named 'minepy'
```

Step 6 (Fixed) — Feature Selection using Mutual Information (MIC Alternative)

`minepy` often fails to install on Kaggle due to compilation issues.
Instead, we use **Mutual Information Regression** from scikit-learn:

- Captures non-linear dependency like MIC
- Stable on Kaggle
- Works well for feature relevance ranking

We compute MI between each feature and the target (`Global_active_power`), then select features with $MI \geq 10\%$ of the max MI (or at least top 5).

Artifact saved:

- `artifacts/selected_features.json`

```
from sklearn.feature_selection import mutual_info_regression

target_col = "Global_active_power"
feature_cols = [c for c in df_hourly.columns if c != target_col]

X_mi = df_hourly[feature_cols].values
y_mi = df_hourly[target_col].values

# Mutual information regression (non-linear)
mi = mutual_info_regression(X_mi, y_mi, random_state=42)

mi_df = pd.DataFrame({"feature": feature_cols, "MI": mi})
mi_df = mi_df.sort_values("MI",
ascending=False).reset_index(drop=True)
```

```

print("Mutual Information scores (high → important):")
display(mi_df)

# Select features using threshold (>= 10% of max) OR at least top 5
max_mi = mi_df["MI"].max()
selected = mi_df[mi_df["MI"] >= 0.10 * max_mi]["feature"].tolist()
if len(selected) < 5:
    selected = mi_df.head(5)["feature"].tolist()

print("Selected features:", selected)

# Save selected features
os.makedirs("/kaggle/working/artifacts", exist_ok=True)
with open("/kaggle/working/artifacts/selected_features.json", "w") as f:
    json.dump(selected, f, indent=2)

print("Saved: /kaggle/working/artifacts/selected_features.json")

Mutual Information scores (high → important):

```

	feature	MI
0	Global_intensity	3.242534
1	Sub_metering_3	0.655366
2	Voltage	0.202940
3	Global_reactive_power	0.173717
4	Sub_metering_2	0.172238
5	Sub_metering_1	0.153720

```

Selected features: ['Global_intensity', 'Sub_metering_3', 'Voltage',
'Global_reactive_power', 'Sub_metering_2']
Saved: /kaggle/working/artifacts/selected_features.json

```

Step 7 — Train/Val/Test Split (Time-based) + Scaling

Important:

- We split by time order (NOT random) to avoid future leakage.
- We fit the scaler ONLY on training data.
- Then transform val/test using the same scaler.

Artifacts saved:

- artifacts/scaler.pkl
- artifacts/config.json

```

target_col = "Global_active_power"

# Use selected features from saved file (ensures consistency)
with open("/kaggle/working/artifacts/selected_features.json", "r") as

```

```

f:
    selected_features = json.load(f)

all_cols = selected_features + [target_col]
data = df_hourly[all_cols].copy()

# Time-based split indices
n = len(data)
train_end = int(0.70 * n)
val_end = int(0.85 * n)

train_df = data.iloc[:train_end]
val_df = data.iloc[train_end:val_end]
test_df = data.iloc[val_end:]

print("□ Splits:")
print("Train:", train_df.shape, "Val:", val_df.shape, "Test:", test_df.shape)

# Fit scaler only on TRAIN
scaler = MinMaxScaler()
scaler.fit(train_df.values)

train_scaled = scaler.transform(train_df.values)
val_scaled = scaler.transform(val_df.values)
test_scaled = scaler.transform(test_df.values)

# Save scaler artifact
joblib.dump(scaler, "/kaggle/working/artifacts/scaler.pkl")
print("□ Saved: /kaggle/working/artifacts/scaler.pkl")

□ Splits:
Train: (24212, 6) Val: (5188, 6) Test: (5189, 6)
□ Saved: /kaggle/working/artifacts/scaler.pkl

```

Step 8 — Sliding Window Dataset Creation

We convert time-series into supervised learning:

- Input: past `lookback` hours of selected features + target history
- Output: next hour `Global_active_power`

Shapes:

- X: (samples, lookback, num_features_total)
- y: (samples, 1)

```

LOOKBACK = 24    # past 24 hours
HORIZON = 1      # predict next 1 hour

# Save config for Flask integration later

```

```

config = {
    "lookback": LOOKBACK,
    "horizon": HORIZON,
    "target_col": target_col,
    "selected_features": selected_features
}
with open("/kaggle/working/artifacts/config.json", "w") as f:
    json.dump(config, f, indent=2)
print("□ Saved: /kaggle/working/artifacts/config.json")

def make_windows(arr, lookback=24, horizon=1, target_index=-1):
    X, y = [], []
    for i in range(len(arr) - lookback - horizon + 1):
        X.append(arr[i:i+lookback])
        y.append(arr[i+lookback+horizon-1, target_index])
    return np.array(X, dtype=np.float32), np.array(y,
dtype=np.float32).reshape(-1, 1)

# Target is last column in our scaled array (because we arranged
# selected_features + target)
target_index = len(all_cols) - 1

X_train, y_train = make_windows(train_scaled, LOOKBACK, HORIZON,
target_index)
X_val, y_val     = make_windows(val_scaled, LOOKBACK, HORIZON,
target_index)
X_test, y_test   = make_windows(test_scaled, LOOKBACK, HORIZON,
target_index)

print("□ Windowed shapes:")
print("X_train:", X_train.shape, "y_train:", y_train.shape)
print("X_val : ", X_val.shape, "y_val : ", y_val.shape)
print("X_test : ", X_test.shape, "y_test : ", y_test.shape)

□ Saved: /kaggle/working/artifacts/config.json
□ Windowed shapes:
X_train: (24188, 24, 6) y_train: (24188, 1)
X_val : (5164, 24, 6) y_val : (5164, 1)
X_test : (5165, 24, 6) y_test : (5165, 1)

```

Step 9 — Build Best Model: CNN + BiLSTM + Self-Attention

Architecture:

- Conv1D (local temporal pattern extraction)
- MaxPool1D
- BiLSTM (long-term dependencies)

- Self-Attention (focus on important time steps)
- Dense (regression output)

We use:

- Adam optimizer
- MSE loss
- Early stopping + model checkpoint

```
class SelfAttention(layers.Layer):
    def __init__(self, **kwargs):
        super().__init__(**kwargs)
        self.att = layers.Attention()

    def call(self, inputs):
        # inputs: (batch, time, features)
        # self-attention: query=key=value=inputs
        context = self.att([inputs, inputs])
        return context

def build_cnn_bilstm_sa(input_shape):
    inp = keras.Input(shape=input_shape)

    x = layers.Conv1D(filters=64, kernel_size=3, padding="same",
activation="relu")(inp)
    x = layers.MaxPooling1D(pool_size=2)(x)
    x = layers.Dropout(0.2)(x)

    x = layers.Bidirectional(layers.LSTM(64, return_sequences=True))(x)
    x = layers.Dropout(0.2)(x)

    x = SelfAttention()(x)
    x = layers.GlobalAveragePooling1D()(x)

    x = layers.Dense(64, activation="relu")(x)
    x = layers.Dropout(0.2)(x)

    out = layers.Dense(1)(x)

    model = keras.Model(inp, out)
    model.compile(optimizer=keras.optimizers.Adam(1e-3), loss="mse",
metrics=["mae"])
    return model

model = build_cnn_bilstm_sa(X_train.shape[1:])
model.summary()

I0000 00:00:1770045205.550725      55 gpu_device.cc:2019] Created
device /job:localhost/replica:0/task:0/device:GPU:0 with 13757 MB
memory: -> device: 0, name: Tesla T4, pci bus id: 0000:00:04.0,
```

```
compute capability: 7.5
I0000 00:00:1770045205.553657      55 gpu_device.cc:2019] Created
device /job:localhost/replica:0/task:0/device:GPU:1 with 13757 MB
memory: -> device: 1, name: Tesla T4, pci bus id: 0000:00:05.0,
compute capability: 7.5
```

Model: "functional"

Layer (type)	Output Shape
Param #	
0 input_layer (InputLayer)	(None, 24, 6)
1,216 conv1d (Conv1D)	(None, 24, 64)
0 max_pooling1d (MaxPooling1D)	(None, 12, 64)
0 dropout (Dropout)	(None, 12, 64)
66,048 bidirectional (Bidirectional)	(None, 12, 128)
0 dropout_1 (Dropout)	(None, 12, 128)
0 self_attention (SelfAttention)	(None, 12, 128)
0 global_average_pooling1d (GlobalAveragePooling1D)	(None, 128)
8,256 dense (Dense)	(None, 64)

```

dropout_2 (Dropout)      | (None, 64)
0 |
dense_1 (Dense)         | (None, 1)
65 |

Total params: 75,585 (295.25 KB)
Trainable params: 75,585 (295.25 KB)
Non-trainable params: 0 (0.00 B)

```

Step 10 — Train CNN–BiLSTM–SelfAttention (Best Practice Training)

We train using:

- EarlyStopping (prevents overfitting)
- ReduceLROnPlateau (improves convergence)
- ModelCheckpoint (saves best model automatically)

Saved:

- `artifacts/model_cnn_bilstm_sa.keras` (best validation loss)

```

# Callbacks (best practice)
ckpt_path = "/kaggle/working/artifacts/model_cnn_bilstm_sa.keras"

callbacks = [
    keras.callbacks.EarlyStopping(
        monitor="val_loss",
        patience=8,
        restore_best_weights=True
    ),
    keras.callbacks.ReduceLROnPlateau(
        monitor="val_loss",
        factor=0.5,
        patience=4,
        min_lr=1e-6,
        verbose=1
    ),
    keras.callbacks.ModelCheckpoint(
        filepath=ckpt_path,
        monitor="val_loss",
        save_best_only=True,

```

```
        verbose=1
    )
]

history = model.fit(
    X_train, y_train,
    validation_data=(X_val, y_val),
    epochs=50,
    batch_size=64,
    callbacks=callbacks,
    verbose=1
)

print("Training complete.")
print("Best model saved at:", ckpt_path)

Epoch 1/50
I0000 00:00:1770045283.415556      191 cuda_dnn.cc:529] Loaded cuDNN
version 91002

375/378 ━━━━━━━━━━ 0s 8ms/step - loss: 0.0172 - mae: 0.0984
Epoch 1: val_loss improved from inf to 0.00881, saving model to
/kaggle/working/artifacts/model_cnn_bilstm_sa.keras
378/378 ━━━━━━━━━━ 10s 10ms/step - loss: 0.0172 - mae:
0.0983 - val_loss: 0.0088 - val_mae: 0.0648 - learning_rate: 0.0010
Epoch 2/50
372/378 ━━━━━━━━━━ 0s 8ms/step - loss: 0.0094 - mae: 0.0676
Epoch 2: val_loss did not improve from 0.00881
378/378 ━━━━━━━━ 3s 8ms/step - loss: 0.0094 - mae: 0.0676
- val_loss: 0.0100 - val_mae: 0.0687 - learning_rate: 0.0010
Epoch 3/50
378/378 ━━━━━━━━━━ 0s 8ms/step - loss: 0.0090 - mae: 0.0662
Epoch 3: val_loss improved from 0.00881 to 0.00851, saving model to
/kaggle/working/artifacts/model_cnn_bilstm_sa.keras
378/378 ━━━━━━━━ 3s 9ms/step - loss: 0.0090 - mae: 0.0662
- val_loss: 0.0085 - val_mae: 0.0646 - learning_rate: 0.0010
Epoch 4/50
372/378 ━━━━━━━━━━ 0s 8ms/step - loss: 0.0087 - mae: 0.0654
Epoch 4: val_loss improved from 0.00851 to 0.00793, saving model to
/kaggle/working/artifacts/model_cnn_bilstm_sa.keras
378/378 ━━━━━━━━ 3s 9ms/step - loss: 0.0087 - mae: 0.0654
- val_loss: 0.0079 - val_mae: 0.0621 - learning_rate: 0.0010
Epoch 5/50
372/378 ━━━━━━━━━━ 0s 8ms/step - loss: 0.0087 - mae: 0.0647
Epoch 5: val_loss improved from 0.00793 to 0.00768, saving model to
/kaggle/working/artifacts/model_cnn_bilstm_sa.keras
378/378 ━━━━━━━━ 3s 9ms/step - loss: 0.0087 - mae: 0.0647
- val_loss: 0.0077 - val_mae: 0.0625 - learning_rate: 0.0010
Epoch 6/50
```

```
372/378 ━━━━━━━━ 0s 7ms/step - loss: 0.0086 - mae: 0.0649
Epoch 6: val_loss improved from 0.00768 to 0.00761, saving model to
/kaggle/working/artifacts/model_cnn_bilstm_sa.keras
378/378 ━━━━━━ 3s 9ms/step - loss: 0.0086 - mae: 0.0648
- val_loss: 0.0076 - val_mae: 0.0642 - learning_rate: 0.0010
Epoch 7/50
377/378 ━━━━━━ 0s 8ms/step - loss: 0.0086 - mae: 0.0648
Epoch 7: val_loss did not improve from 0.00761
378/378 ━━━━━━ 3s 8ms/step - loss: 0.0086 - mae: 0.0648
- val_loss: 0.0080 - val_mae: 0.0637 - learning_rate: 0.0010
Epoch 8/50
372/378 ━━━━━━ 0s 8ms/step - loss: 0.0084 - mae: 0.0644
Epoch 8: val_loss did not improve from 0.00761
378/378 ━━━━━━ 3s 8ms/step - loss: 0.0084 - mae: 0.0644
- val_loss: 0.0077 - val_mae: 0.0619 - learning_rate: 0.0010
Epoch 9/50
372/378 ━━━━━━ 0s 8ms/step - loss: 0.0084 - mae: 0.0636
Epoch 9: val_loss improved from 0.00761 to 0.00749, saving model to
/kaggle/working/artifacts/model_cnn_bilstm_sa.keras
378/378 ━━━━━━ 3s 9ms/step - loss: 0.0084 - mae: 0.0636
- val_loss: 0.0075 - val_mae: 0.0628 - learning_rate: 0.0010
Epoch 10/50
375/378 ━━━━━━ 0s 8ms/step - loss: 0.0081 - mae: 0.0628
Epoch 10: val_loss did not improve from 0.00749
378/378 ━━━━━━ 3s 9ms/step - loss: 0.0081 - mae: 0.0628
- val_loss: 0.0078 - val_mae: 0.0618 - learning_rate: 0.0010
Epoch 11/50
372/378 ━━━━━━ 0s 8ms/step - loss: 0.0083 - mae: 0.0638
Epoch 11: val_loss did not improve from 0.00749
378/378 ━━━━━━ 3s 8ms/step - loss: 0.0083 - mae: 0.0638
- val_loss: 0.0075 - val_mae: 0.0595 - learning_rate: 0.0010
Epoch 12/50
372/378 ━━━━━━ 0s 8ms/step - loss: 0.0080 - mae: 0.0621
Epoch 12: val_loss did not improve from 0.00749
378/378 ━━━━━━ 3s 8ms/step - loss: 0.0080 - mae: 0.0621
- val_loss: 0.0078 - val_mae: 0.0614 - learning_rate: 0.0010
Epoch 13/50
375/378 ━━━━━━ 0s 8ms/step - loss: 0.0081 - mae: 0.0627
Epoch 13: val_loss improved from 0.00749 to 0.00732, saving model
to /kaggle/working/artifacts/model_cnn_bilstm_sa.keras
378/378 ━━━━━━ 3s 9ms/step - loss: 0.0081 - mae: 0.0627
- val_loss: 0.0073 - val_mae: 0.0606 - learning_rate: 0.0010
Epoch 14/50
372/378 ━━━━━━ 0s 8ms/step - loss: 0.0083 - mae: 0.0633
Epoch 14: val_loss improved from 0.00732 to 0.00724, saving model
to /kaggle/working/artifacts/model_cnn_bilstm_sa.keras
378/378 ━━━━━━ 3s 9ms/step - loss: 0.0083 - mae: 0.0632
- val_loss: 0.0072 - val_mae: 0.0611 - learning_rate: 0.0010
Epoch 15/50
```

```
372/378 ━━━━━━━━ 0s 7ms/step - loss: 0.0085 - mae: 0.0638
Epoch 15: val_loss did not improve from 0.00724
378/378 ━━━━━━ 3s 8ms/step - loss: 0.0085 - mae: 0.0638
- val_loss: 0.0075 - val_mae: 0.0597 - learning_rate: 0.0010
Epoch 16/50
375/378 ━━━━━━ 0s 8ms/step - loss: 0.0081 - mae: 0.0624
Epoch 16: val_loss did not improve from 0.00724
378/378 ━━━━━━ 3s 8ms/step - loss: 0.0081 - mae: 0.0624
- val_loss: 0.0074 - val_mae: 0.0609 - learning_rate: 0.0010
Epoch 17/50
372/378 ━━━━━━ 0s 7ms/step - loss: 0.0082 - mae: 0.0629
Epoch 17: val_loss improved from 0.00724 to 0.00718, saving model
to /kaggle/working/artifacts/model_cnn_bilstm_sa.keras
378/378 ━━━━━━ 3s 8ms/step - loss: 0.0082 - mae: 0.0629
- val_loss: 0.0072 - val_mae: 0.0605 - learning_rate: 0.0010
Epoch 18/50
372/378 ━━━━━━ 0s 7ms/step - loss: 0.0082 - mae: 0.0631
Epoch 18: val_loss did not improve from 0.00718
378/378 ━━━━━━ 3s 8ms/step - loss: 0.0082 - mae: 0.0630
- val_loss: 0.0072 - val_mae: 0.0623 - learning_rate: 0.0010
Epoch 19/50
376/378 ━━━━━━ 0s 7ms/step - loss: 0.0080 - mae: 0.0624
Epoch 19: val_loss improved from 0.00718 to 0.00718, saving model
to /kaggle/working/artifacts/model_cnn_bilstm_sa.keras
378/378 ━━━━━━ 3s 8ms/step - loss: 0.0080 - mae: 0.0624
- val_loss: 0.0072 - val_mae: 0.0592 - learning_rate: 0.0010
Epoch 20/50
372/378 ━━━━━━ 0s 7ms/step - loss: 0.0078 - mae: 0.0616
Epoch 20: val_loss did not improve from 0.00718
378/378 ━━━━━━ 3s 8ms/step - loss: 0.0078 - mae: 0.0616
- val_loss: 0.0082 - val_mae: 0.0622 - learning_rate: 0.0010
Epoch 21/50
372/378 ━━━━━━ 0s 7ms/step - loss: 0.0079 - mae: 0.0616
Epoch 21: ReduceLROnPlateau reducing learning rate to
0.0005000000237487257.

Epoch 21: val_loss improved from 0.00718 to 0.00712, saving model
to /kaggle/working/artifacts/model_cnn_bilstm_sa.keras
378/378 ━━━━━━ 3s 8ms/step - loss: 0.0079 - mae: 0.0615
- val_loss: 0.0071 - val_mae: 0.0612 - learning_rate: 0.0010
Epoch 22/50
376/378 ━━━━━━ 0s 8ms/step - loss: 0.0076 - mae: 0.0602
Epoch 22: val_loss did not improve from 0.00712
378/378 ━━━━━━ 3s 8ms/step - loss: 0.0076 - mae: 0.0602
- val_loss: 0.0072 - val_mae: 0.0594 - learning_rate: 5.0000e-04
Epoch 23/50
372/378 ━━━━━━ 0s 7ms/step - loss: 0.0075 - mae: 0.0606
Epoch 23: val_loss did not improve from 0.00712
378/378 ━━━━━━ 3s 8ms/step - loss: 0.0075 - mae: 0.0606
```

```
- val_loss: 0.0072 - val_mae: 0.0593 - learning_rate: 5.0000e-04
Epoch 24/50
372/378 ━━━━━━━━ 0s 7ms/step - loss: 0.0077 - mae: 0.0605
Epoch 24: val_loss did not improve from 0.00712
378/378 ━━━━━━ 3s 8ms/step - loss: 0.0077 - mae: 0.0605
- val_loss: 0.0072 - val_mae: 0.0594 - learning_rate: 5.0000e-04
Epoch 25/50
374/378 ━━━━━━ 0s 8ms/step - loss: 0.0076 - mae: 0.0603
Epoch 25: ReduceLROnPlateau reducing learning rate to
0.0002500000118743628.

Epoch 25: val_loss improved from 0.00712 to 0.00709, saving model
to /kaggle/working/artifacts/model_cnn_bilstm_sa.keras
378/378 ━━━━━━ 3s 9ms/step - loss: 0.0076 - mae: 0.0603
- val_loss: 0.0071 - val_mae: 0.0588 - learning_rate: 5.0000e-04
Epoch 26/50
372/378 ━━━━━━ 0s 8ms/step - loss: 0.0076 - mae: 0.0598
Epoch 26: val_loss improved from 0.00709 to 0.00696, saving model
to /kaggle/working/artifacts/model_cnn_bilstm_sa.keras
378/378 ━━━━━━ 3s 9ms/step - loss: 0.0076 - mae: 0.0598
- val_loss: 0.0070 - val_mae: 0.0593 - learning_rate: 2.5000e-04
Epoch 27/50
372/378 ━━━━━━ 0s 8ms/step - loss: 0.0075 - mae: 0.0597
Epoch 27: val_loss did not improve from 0.00696
378/378 ━━━━━━ 3s 8ms/step - loss: 0.0075 - mae: 0.0597
- val_loss: 0.0070 - val_mae: 0.0587 - learning_rate: 2.5000e-04
Epoch 28/50
376/378 ━━━━━━ 0s 8ms/step - loss: 0.0075 - mae: 0.0598
Epoch 28: val_loss did not improve from 0.00696
378/378 ━━━━━━ 3s 9ms/step - loss: 0.0075 - mae: 0.0598
- val_loss: 0.0070 - val_mae: 0.0594 - learning_rate: 2.5000e-04
Epoch 29/50
372/378 ━━━━━━ 0s 8ms/step - loss: 0.0075 - mae: 0.0600
Epoch 29: val_loss did not improve from 0.00696
378/378 ━━━━━━ 3s 8ms/step - loss: 0.0075 - mae: 0.0600
- val_loss: 0.0071 - val_mae: 0.0594 - learning_rate: 2.5000e-04
Epoch 30/50
372/378 ━━━━━━ 0s 8ms/step - loss: 0.0075 - mae: 0.0596
Epoch 30: ReduceLROnPlateau reducing learning rate to
0.0001250000059371814.

Epoch 30: val_loss did not improve from 0.00696
378/378 ━━━━━━ 3s 8ms/step - loss: 0.0075 - mae: 0.0596
- val_loss: 0.0071 - val_mae: 0.0589 - learning_rate: 2.5000e-04
Epoch 31/50
376/378 ━━━━━━ 0s 8ms/step - loss: 0.0074 - mae: 0.0593
Epoch 31: val_loss improved from 0.00696 to 0.00695, saving model
to /kaggle/working/artifacts/model_cnn_bilstm_sa.keras
378/378 ━━━━━━ 3s 9ms/step - loss: 0.0074 - mae: 0.0593
```

```
- val_loss: 0.0069 - val_mae: 0.0590 - learning_rate: 1.2500e-04
Epoch 32/50
372/378 ----- 0s 7ms/step - loss: 0.0074 - mae: 0.0594
Epoch 32: val_loss did not improve from 0.00695
378/378 ----- 3s 8ms/step - loss: 0.0074 - mae: 0.0594
- val_loss: 0.0071 - val_mae: 0.0593 - learning_rate: 1.2500e-04
Epoch 33/50
372/378 ----- 0s 7ms/step - loss: 0.0074 - mae: 0.0593
Epoch 33: val_loss did not improve from 0.00695
378/378 ----- 3s 8ms/step - loss: 0.0074 - mae: 0.0593
- val_loss: 0.0072 - val_mae: 0.0592 - learning_rate: 1.2500e-04
Epoch 34/50
372/378 ----- 0s 8ms/step - loss: 0.0074 - mae: 0.0589
Epoch 34: ReduceLROnPlateau reducing learning rate to
6.25000029685907e-05.

Epoch 34: val_loss did not improve from 0.00695
378/378 ----- 3s 8ms/step - loss: 0.0074 - mae: 0.0589
- val_loss: 0.0071 - val_mae: 0.0593 - learning_rate: 1.2500e-04
Epoch 35/50
372/378 ----- 0s 8ms/step - loss: 0.0073 - mae: 0.0588
Epoch 35: val_loss did not improve from 0.00695
378/378 ----- 3s 8ms/step - loss: 0.0073 - mae: 0.0588
- val_loss: 0.0070 - val_mae: 0.0585 - learning_rate: 6.2500e-05
Epoch 36/50
372/378 ----- 0s 7ms/step - loss: 0.0075 - mae: 0.0593
Epoch 36: val_loss did not improve from 0.00695
378/378 ----- 3s 8ms/step - loss: 0.0075 - mae: 0.0593
- val_loss: 0.0070 - val_mae: 0.0590 - learning_rate: 6.2500e-05
Epoch 37/50
372/378 ----- 0s 7ms/step - loss: 0.0074 - mae: 0.0588
Epoch 37: val_loss did not improve from 0.00695
378/378 ----- 3s 8ms/step - loss: 0.0074 - mae: 0.0588
- val_loss: 0.0071 - val_mae: 0.0590 - learning_rate: 6.2500e-05
Epoch 38/50
377/378 ----- 0s 8ms/step - loss: 0.0073 - mae: 0.0588
Epoch 38: ReduceLROnPlateau reducing learning rate to
3.125000148429535e-05.

Epoch 38: val_loss did not improve from 0.00695
378/378 ----- 3s 8ms/step - loss: 0.0073 - mae: 0.0588
- val_loss: 0.0070 - val_mae: 0.0590 - learning_rate: 6.2500e-05
Epoch 39/50
372/378 ----- 0s 8ms/step - loss: 0.0074 - mae: 0.0592
Epoch 39: val_loss did not improve from 0.00695
378/378 ----- 3s 8ms/step - loss: 0.0074 - mae: 0.0592
- val_loss: 0.0070 - val_mae: 0.0588 - learning_rate: 3.1250e-05
[] Training complete.
```

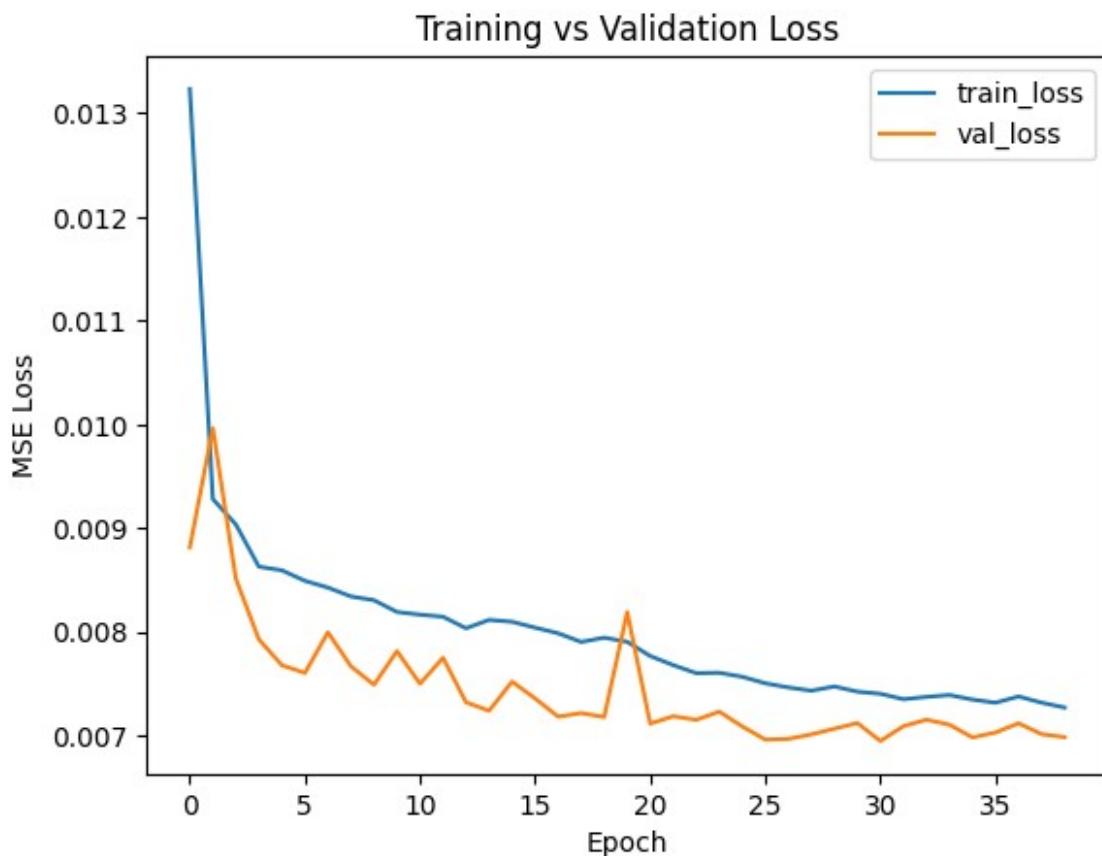
```

□ Best model saved at:  

/kaggle/working/artifacts/model_cnn_bilstm_sa.keras

plt.figure()
plt.plot(history.history["loss"], label="train_loss")
plt.plot(history.history["val_loss"], label="val_loss")
plt.xlabel("Epoch")
plt.ylabel("MSE Loss")
plt.legend()
plt.title("Training vs Validation Loss")
plt.show()

```



Step 11 — Evaluation (RMSE, MAE, R²)

We evaluate on the test set using:

- RMSE (lower is better)
- MAE (lower is better)
- R² (closer to 1 is better)

We also plot:

- Actual vs Predicted (first 300 points)

```

# Load best saved model (ensures we evaluate the best checkpoint)
best_model = keras.models.load_model(
    "/kaggle/working/artifacts/model_cnn_bilstm_sa.keras",
    custom_objects={"SelfAttention": SelfAttention}
)

# Predict
y_pred = best_model.predict(X_test, verbose=0)

# Metrics
rmse = math.sqrt(mean_squared_error(y_test, y_pred))
mae = mean_absolute_error(y_test, y_pred)
r2 = r2_score(y_test, y_pred)

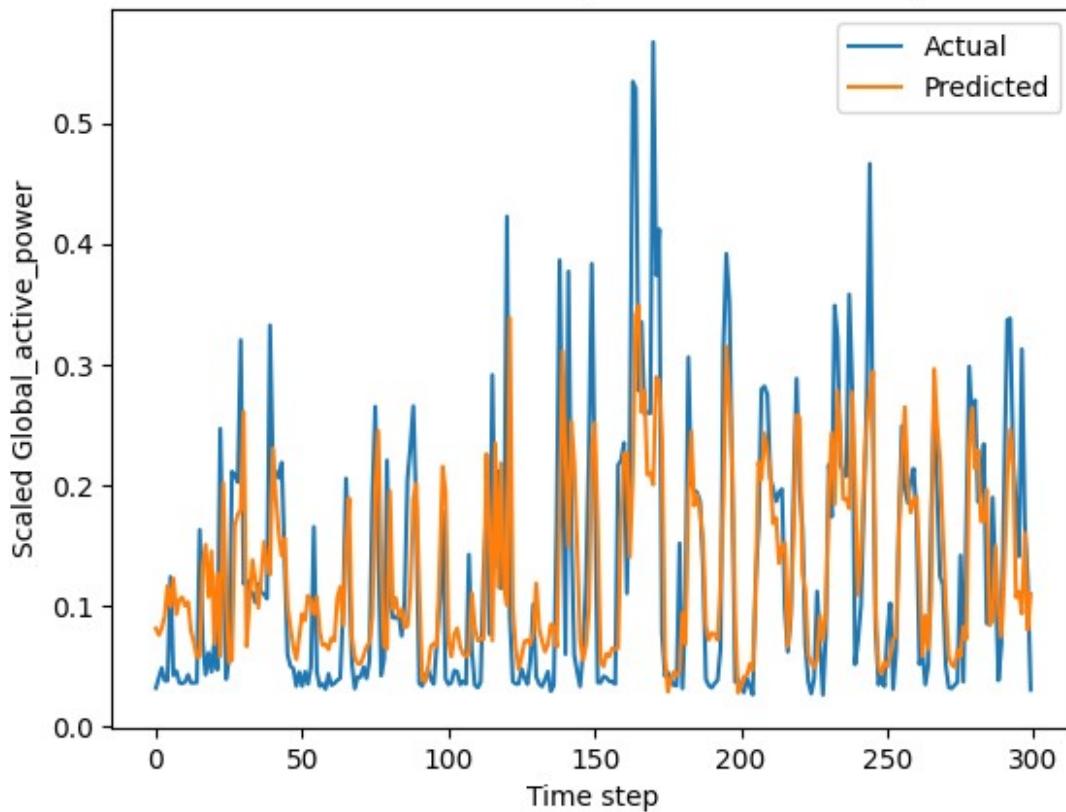
print("□ Test Metrics (Scaled target space):")
print("RMSE:", rmse)
print("MAE :", mae)
print("R2 :", r2)

□ Test Metrics (Scaled target space):
RMSE: 0.07225319394597163
MAE : 0.05097724869847298
R2 : 0.5625068545341492

plt.figure()
plt.plot(y_test[:300], label="Actual")
plt.plot(y_pred[:300], label="Predicted")
plt.title("Actual vs Predicted (First 300 Test Points)")
plt.xlabel("Time step")
plt.ylabel("Scaled Global_active_power")
plt.legend()
plt.show()

```

Actual vs Predicted (First 300 Test Points)



Step 12 — Convert Predictions Back to Original Units + Save Artifacts

Because we scaled all columns together, inverse-transforming the target alone requires:

1. Create a dummy array with the same number of columns
2. Put predicted target into the last column
3. Use `scaler.inverse_transform()`
4. Extract the last column as the true kW values

We also save:

- `metrics.json`
- model + scaler + config + selected_features All inside `/kaggle/working/artifacts/` so you can download them easily.

```
# Helper: inverse transform only the target (last column)
def inverse_target(scaler, y_scaled, num_cols, target_index):
    dummy = np.zeros((len(y_scaled), num_cols), dtype=np.float32)
    dummy[:, target_index] = y_scaled.reshape(-1)
    inv = scaler.inverse_transform(dummy)
    return inv[:, target_index]
```

```

num_cols = len(all_cols)
t_idx = target_index

# Inverse transform y_test and y_pred to original kW units
y_test_kw = inverse_target(scaler, y_test, num_cols, t_idx)
y_pred_kw = inverse_target(scaler, y_pred, num_cols, t_idx)

# Metrics in original units
rmse_kw = math.sqrt(mean_squared_error(y_test_kw, y_pred_kw))
mae_kw = mean_absolute_error(y_test_kw, y_pred_kw)
r2_kw = r2_score(y_test_kw, y_pred_kw)

print("□ Test Metrics (Original kW units):")
print("RMSE (kW):", rmse_kw)
print("MAE (kW):", mae_kw)
print("R2      :", r2_kw)

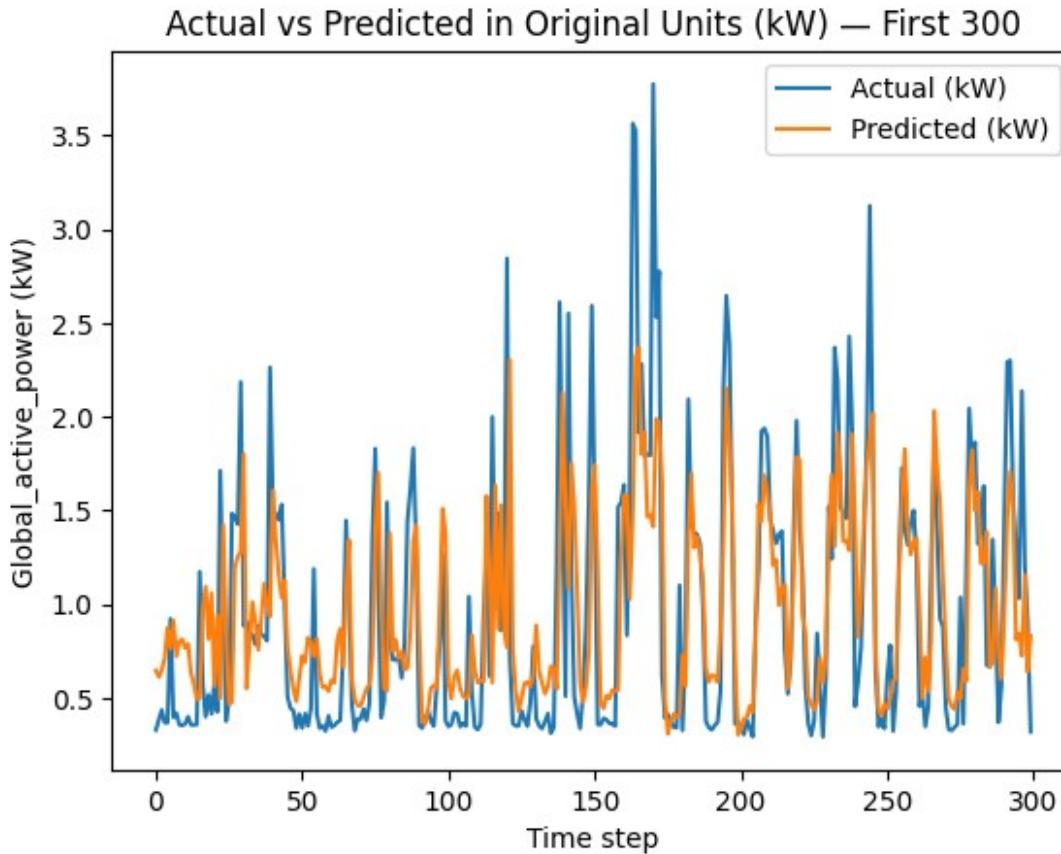
# Save metrics
metrics = {
    "scaled": {"rmse": float(rmse), "mae": float(mae), "r2": float(r2)},
    "original_kw": {"rmse": float(rmse_kw), "mae": float(mae_kw), "r2": float(r2_kw)},
    "selected_features": selected_features,
    "lookback": LOOKBACK,
    "horizon": HORIZON
}
with open("/kaggle/working/artifacts/metrics_cnn_bilstm_sa.json", "w") as f:
    json.dump(metrics, f, indent=2)

print("□ Saved: /kaggle/working/artifacts/metrics_cnn_bilstm_sa.json")

□ Test Metrics (Original kW units):
RMSE (kW): 0.4650600955697453
MAE (kW): 0.3281167149543762
R2      : 0.5625068545341492
□ Saved: /kaggle/working/artifacts/metrics_cnn_bilstm_sa.json

plt.figure()
plt.plot(y_test_kw[:300], label="Actual (kW)")
plt.plot(y_pred_kw[:300], label="Predicted (kW)")
plt.title("Actual vs Predicted in Original Units (kW) – First 300")
plt.xlabel("Time step")
plt.ylabel("Global_active_power (kW)")
plt.legend()
plt.show()

```



Step 13 — Build CNN + BiGRU + Self-Attention (Faster Variant)

GRU is lighter than LSTM (fewer gates), often faster and sometimes performs better. We keep everything same:

- Conv1D → MaxPool → Dropout
- BiGRU (return_sequences=True)
- Self-Attention
- GlobalAveragePooling
- Dense → Output(1)

```
def build_cnn_bigru_sa(input_shape):
    inp = keras.Input(shape=input_shape)

    x = layers.Conv1D(filters=64, kernel_size=3, padding="same",
activation="relu")(inp)
    x = layers.MaxPooling1D(pool_size=2)(x)
    x = layers.Dropout(0.2)(x)

    x = layers.Bidirectional(layers.GRU(64, return_sequences=True))(x)
    x = layers.Dropout(0.2)(x)
```

```

x = SelfAttention()(x)
x = layers.GlobalAveragePooling1D()(x)

x = layers.Dense(64, activation="relu")(x)
x = layers.Dropout(0.2)(x)

out = layers.Dense(1)(x)

model = keras.Model(inp, out)
model.compile(optimizer=keras.optimizers.Adam(1e-3), loss="mse",
metrics=["mae"])
return model

model_bigru = build_cnn_bigru_sa(X_train.shape[1:])
model_bigru.summary()

```

Model: "functional_1"

Layer (type)	Output Shape
Param #	
input_layer_1 (InputLayer)	(None, 24, 6)
0	
conv1d_1 (Conv1D)	(None, 24, 64)
1,216	
max_pooling1d_1 (MaxPooling1D)	(None, 12, 64)
0	
dropout_3 (Dropout)	(None, 12, 64)
0	
bidirectional_1 (Bidirectional)	(None, 12, 128)
49,920	
dropout_4 (Dropout)	(None, 12, 128)
0	
self_attention_1	(None, 12, 128)

```

0 | (SelfAttention)
  |
  |
  | global_average_pooling1d_1 | (None, 128)
0 | (GlobalAveragePooling1D)
  |
  |
  | dense_2 (Dense) | (None, 64)
8,256 |
  |
  | dropout_5 (Dropout) | (None, 64)
0 |
  |
  | dense_3 (Dense) | (None, 1)
65 |
  |
  |
Total params: 59,457 (232.25 KB)
Trainable params: 59,457 (232.25 KB)
Non-trainable params: 0 (0.00 B)

```

Step 14 — Train CNN–BiGRU–SelfAttention

We reuse the same training setup:

- EarlyStopping
- ReduceLROnPlateau
- ModelCheckpoint

Saved:

- artifacts/model_cnn_bigru_sa.keras

```

ckpt_path_bigru = "/kaggle/working/artifacts/model_cnn_bigru_sa.keras"

callbacks_bigru = [
    keras.callbacks.EarlyStopping(
        monitor="val_loss",
        patience=8,
        restore_best_weights=True
),

```

```

        keras.callbacks.ReduceLROnPlateau(
            monitor="val_loss",
            factor=0.5,
            patience=4,
            min_lr=1e-6,
            verbose=1
        ),
        keras.callbacks.ModelCheckpoint(
            filepath=ckpt_path_bigru,
            monitor="val_loss",
            save_best_only=True,
            verbose=1
        )
    )

history_bigru = model_bigru.fit(
    X_train, y_train,
    validation_data=(X_val, y_val),
    epochs=50,
    batch_size=64,
    callbacks=callbacks_bigru,
    verbose=1
)

print("[] BiGRU Training complete.")
print("[] Best BiGRU model saved at:", ckpt_path_bigru)

Epoch 1/50
372/378 ━━━━━━━━ 0s 7ms/step - loss: 0.0174 - mae: 0.1004
Epoch 1: val_loss improved from inf to 0.00886, saving model to
/kaggle/working/artifacts/model_cnn_bigru_sa.keras
378/378 ━━━━━━ 6s 9ms/step - loss: 0.0173 - mae: 0.1002
- val_loss: 0.0089 - val_mae: 0.0715 - learning_rate: 0.0010
Epoch 2/50
375/378 ━━━━━━ 0s 8ms/step - loss: 0.0102 - mae: 0.0716
Epoch 2: val_loss improved from 0.00886 to 0.00827, saving model to
/kaggle/working/artifacts/model_cnn_bigru_sa.keras
378/378 ━━━━ 3s 9ms/step - loss: 0.0102 - mae: 0.0716
- val_loss: 0.0083 - val_mae: 0.0645 - learning_rate: 0.0010
Epoch 3/50
372/378 ━━━━ 0s 7ms/step - loss: 0.0093 - mae: 0.0676
Epoch 3: val_loss did not improve from 0.00827
378/378 ━━━━ 3s 8ms/step - loss: 0.0093 - mae: 0.0676
- val_loss: 0.0083 - val_mae: 0.0652 - learning_rate: 0.0010
Epoch 4/50
372/378 ━━━━ 0s 7ms/step - loss: 0.0094 - mae: 0.0687
Epoch 4: val_loss did not improve from 0.00827
378/378 ━━━━ 3s 8ms/step - loss: 0.0094 - mae: 0.0687
- val_loss: 0.0097 - val_mae: 0.0673 - learning_rate: 0.0010
Epoch 5/50

```

```
377/378 ----- 0s 7ms/step - loss: 0.0090 - mae: 0.0663
Epoch 5: val_loss did not improve from 0.00827
378/378 ----- 3s 8ms/step - loss: 0.0090 - mae: 0.0663
- val_loss: 0.0092 - val_mae: 0.0655 - learning_rate: 0.0010
Epoch 6/50
373/378 ----- 0s 7ms/step - loss: 0.0088 - mae: 0.0654
Epoch 6: val_loss improved from 0.00827 to 0.00781, saving model to
/kaggle/working/artifacts/model_cnn_bigru_sa.keras
378/378 ----- 3s 8ms/step - loss: 0.0088 - mae: 0.0654
- val_loss: 0.0078 - val_mae: 0.0634 - learning_rate: 0.0010
Epoch 7/50
372/378 ----- 0s 7ms/step - loss: 0.0087 - mae: 0.0652
Epoch 7: val_loss improved from 0.00781 to 0.00769, saving model to
/kaggle/working/artifacts/model_cnn_bigru_sa.keras
378/378 ----- 3s 8ms/step - loss: 0.0087 - mae: 0.0652
- val_loss: 0.0077 - val_mae: 0.0624 - learning_rate: 0.0010
Epoch 8/50
375/378 ----- 0s 8ms/step - loss: 0.0087 - mae: 0.0650
Epoch 8: val_loss improved from 0.00769 to 0.00750, saving model to
/kaggle/working/artifacts/model_cnn_bigru_sa.keras
378/378 ----- 3s 9ms/step - loss: 0.0087 - mae: 0.0650
- val_loss: 0.0075 - val_mae: 0.0623 - learning_rate: 0.0010
Epoch 9/50
372/378 ----- 0s 7ms/step - loss: 0.0086 - mae: 0.0653
Epoch 9: val_loss did not improve from 0.00750
378/378 ----- 3s 8ms/step - loss: 0.0086 - mae: 0.0653
- val_loss: 0.0091 - val_mae: 0.0657 - learning_rate: 0.0010
Epoch 10/50
372/378 ----- 0s 7ms/step - loss: 0.0084 - mae: 0.0638
Epoch 10: val_loss did not improve from 0.00750
378/378 ----- 3s 8ms/step - loss: 0.0084 - mae: 0.0639
- val_loss: 0.0077 - val_mae: 0.0610 - learning_rate: 0.0010
Epoch 11/50
377/378 ----- 0s 8ms/step - loss: 0.0083 - mae: 0.0640
Epoch 11: val_loss improved from 0.00750 to 0.00740, saving model
to /kaggle/working/artifacts/model_cnn_bigru_sa.keras
378/378 ----- 3s 9ms/step - loss: 0.0083 - mae: 0.0640
- val_loss: 0.0074 - val_mae: 0.0627 - learning_rate: 0.0010
Epoch 12/50
372/378 ----- 0s 8ms/step - loss: 0.0086 - mae: 0.0643
Epoch 12: ReduceLROnPlateau reducing learning rate to
0.000500000237487257.

Epoch 12: val_loss did not improve from 0.00740
378/378 ----- 3s 8ms/step - loss: 0.0086 - mae: 0.0643
- val_loss: 0.0079 - val_mae: 0.0615 - learning_rate: 0.0010
Epoch 13/50
372/378 ----- 0s 7ms/step - loss: 0.0082 - mae: 0.0627
Epoch 13: val_loss improved from 0.00740 to 0.00732, saving model
```

```
to /kaggle/working/artifacts/model_cnn_bigru_sa.keras
378/378 ————— 3s 8ms/step - loss: 0.0082 - mae: 0.0627
- val_loss: 0.0073 - val_mae: 0.0604 - learning_rate: 5.0000e-04
Epoch 14/50
377/378 ————— 0s 8ms/step - loss: 0.0080 - mae: 0.0627
Epoch 14: val_loss did not improve from 0.00732
378/378 ————— 3s 9ms/step - loss: 0.0080 - mae: 0.0627
- val_loss: 0.0073 - val_mae: 0.0611 - learning_rate: 5.0000e-04
Epoch 15/50
372/378 ————— 0s 7ms/step - loss: 0.0080 - mae: 0.0625
Epoch 15: val_loss did not improve from 0.00732
378/378 ————— 3s 8ms/step - loss: 0.0080 - mae: 0.0625
- val_loss: 0.0073 - val_mae: 0.0608 - learning_rate: 5.0000e-04
Epoch 16/50
372/378 ————— 0s 7ms/step - loss: 0.0081 - mae: 0.0633
Epoch 16: val_loss did not improve from 0.00732
378/378 ————— 3s 8ms/step - loss: 0.0081 - mae: 0.0633
- val_loss: 0.0074 - val_mae: 0.0619 - learning_rate: 5.0000e-04
Epoch 17/50
372/378 ————— 0s 8ms/step - loss: 0.0080 - mae: 0.0619
Epoch 17: ReduceLROnPlateau reducing learning rate to
0.0002500000118743628.

Epoch 17: val_loss did not improve from 0.00732
378/378 ————— 3s 9ms/step - loss: 0.0080 - mae: 0.0619
- val_loss: 0.0075 - val_mae: 0.0602 - learning_rate: 5.0000e-04
Epoch 18/50
378/378 ————— 0s 8ms/step - loss: 0.0078 - mae: 0.0610
Epoch 18: val_loss improved from 0.00732 to 0.00722, saving model
to /kaggle/working/artifacts/model_cnn_bigru_sa.keras
378/378 ————— 3s 9ms/step - loss: 0.0078 - mae: 0.0610
- val_loss: 0.0072 - val_mae: 0.0598 - learning_rate: 2.5000e-04
Epoch 19/50
372/378 ————— 0s 7ms/step - loss: 0.0080 - mae: 0.0619
Epoch 19: val_loss did not improve from 0.00722
378/378 ————— 3s 8ms/step - loss: 0.0080 - mae: 0.0619
- val_loss: 0.0077 - val_mae: 0.0607 - learning_rate: 2.5000e-04
Epoch 20/50
373/378 ————— 0s 7ms/step - loss: 0.0077 - mae: 0.0610
Epoch 20: val_loss improved from 0.00722 to 0.00715, saving model
to /kaggle/working/artifacts/model_cnn_bigru_sa.keras
378/378 ————— 3s 8ms/step - loss: 0.0077 - mae: 0.0610
- val_loss: 0.0071 - val_mae: 0.0600 - learning_rate: 2.5000e-04
Epoch 21/50
373/378 ————— 0s 7ms/step - loss: 0.0077 - mae: 0.0605
Epoch 21: val_loss did not improve from 0.00715
378/378 ————— 3s 8ms/step - loss: 0.0077 - mae: 0.0605
- val_loss: 0.0072 - val_mae: 0.0599 - learning_rate: 2.5000e-04
Epoch 22/50
```

```
374/378 ----- 0s 7ms/step - loss: 0.0077 - mae: 0.0608
Epoch 22: ReduceLROnPlateau reducing learning rate to
0.0001250000059371814.

Epoch 22: val_loss did not improve from 0.00715
378/378 ----- 3s 8ms/step - loss: 0.0077 - mae: 0.0608
- val_loss: 0.0072 - val_mae: 0.0603 - learning_rate: 2.5000e-04
Epoch 23/50
374/378 ----- 0s 7ms/step - loss: 0.0076 - mae: 0.0606
Epoch 23: val_loss improved from 0.00715 to 0.00709, saving model
to /kaggle/working/artifacts/model_cnn_bigru_sa.keras
378/378 ----- 3s 8ms/step - loss: 0.0076 - mae: 0.0606
- val_loss: 0.0071 - val_mae: 0.0599 - learning_rate: 1.2500e-04
Epoch 24/50
376/378 ----- 0s 7ms/step - loss: 0.0078 - mae: 0.0612
Epoch 24: val_loss did not improve from 0.00709
378/378 ----- 3s 8ms/step - loss: 0.0078 - mae: 0.0612
- val_loss: 0.0073 - val_mae: 0.0598 - learning_rate: 1.2500e-04
Epoch 25/50
372/378 ----- 0s 7ms/step - loss: 0.0076 - mae: 0.0602
Epoch 25: val_loss did not improve from 0.00709
378/378 ----- 3s 8ms/step - loss: 0.0076 - mae: 0.0603
- val_loss: 0.0071 - val_mae: 0.0594 - learning_rate: 1.2500e-04
Epoch 26/50
372/378 ----- 0s 7ms/step - loss: 0.0080 - mae: 0.0616
Epoch 26: val_loss did not improve from 0.00709
378/378 ----- 3s 8ms/step - loss: 0.0080 - mae: 0.0616
- val_loss: 0.0072 - val_mae: 0.0596 - learning_rate: 1.2500e-04
Epoch 27/50
376/378 ----- 0s 8ms/step - loss: 0.0080 - mae: 0.0615
Epoch 27: ReduceLROnPlateau reducing learning rate to
6.25000029685907e-05.

Epoch 27: val_loss did not improve from 0.00709
378/378 ----- 3s 8ms/step - loss: 0.0080 - mae: 0.0615
- val_loss: 0.0072 - val_mae: 0.0595 - learning_rate: 1.2500e-04
Epoch 28/50
372/378 ----- 0s 7ms/step - loss: 0.0076 - mae: 0.0604
Epoch 28: val_loss did not improve from 0.00709
378/378 ----- 3s 8ms/step - loss: 0.0076 - mae: 0.0604
- val_loss: 0.0072 - val_mae: 0.0595 - learning_rate: 6.2500e-05
Epoch 29/50
372/378 ----- 0s 7ms/step - loss: 0.0074 - mae: 0.0596
Epoch 29: val_loss did not improve from 0.00709
378/378 ----- 3s 8ms/step - loss: 0.0074 - mae: 0.0596
- val_loss: 0.0071 - val_mae: 0.0594 - learning_rate: 6.2500e-05
Epoch 30/50
375/378 ----- 0s 8ms/step - loss: 0.0076 - mae: 0.0604
Epoch 30: val_loss did not improve from 0.00709
```

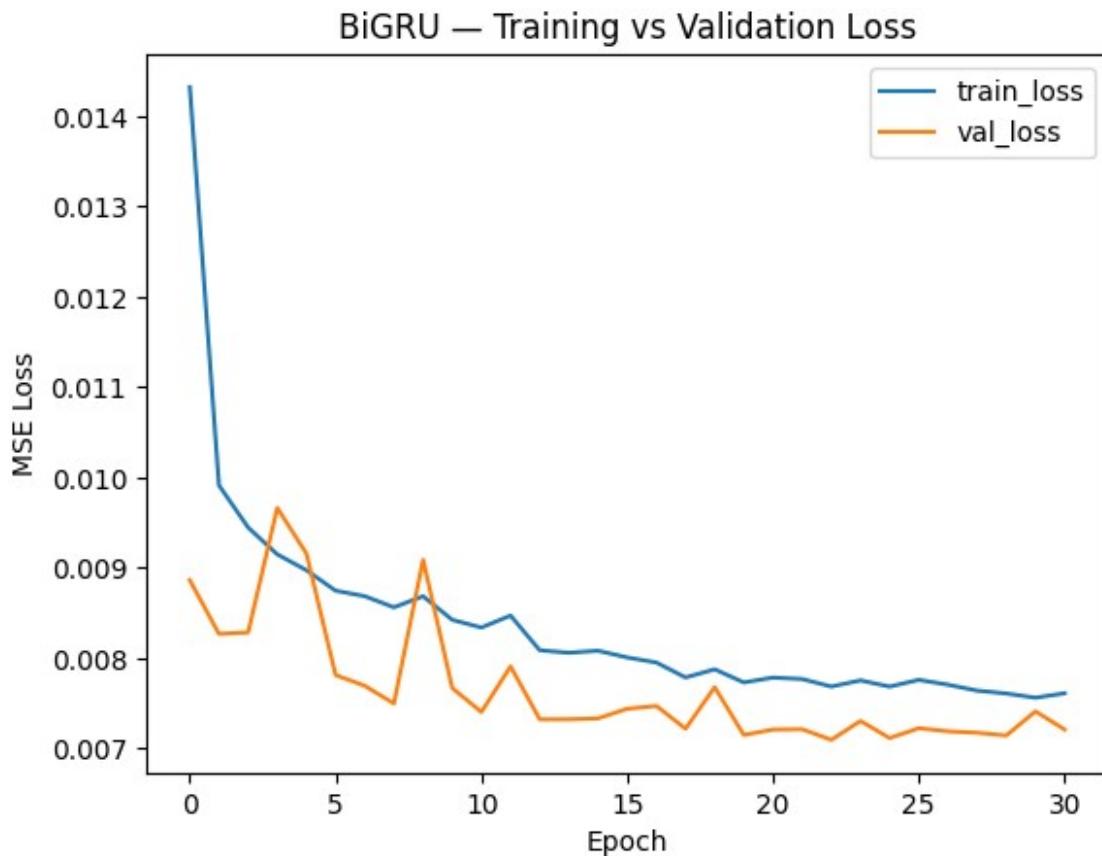
```

378/378 ━━━━━━━━━━ 3s 9ms/step - loss: 0.0076 - mae: 0.0604
- val_loss: 0.0074 - val_mae: 0.0605 - learning_rate: 6.2500e-05
Epoch 31/50
372/378 ━━━━━━━━ 0s 8ms/step - loss: 0.0077 - mae: 0.0608
Epoch 31: ReduceLROnPlateau reducing learning rate to
3.125000148429535e-05.

Epoch 31: val_loss did not improve from 0.00709
378/378 ━━━━━━━━ 3s 8ms/step - loss: 0.0077 - mae: 0.0608
- val_loss: 0.0072 - val_mae: 0.0597 - learning_rate: 6.2500e-05
□ BiGRU Training complete.
□ Best BiGRU model saved at:
/kaggle/working/artifacts/model_cnn_bigru_sa.keras

plt.figure()
plt.plot(history_bigru.history["loss"], label="train_loss")
plt.plot(history_bigru.history["val_loss"], label="val_loss")
plt.xlabel("Epoch")
plt.ylabel("MSE Loss")
plt.legend()
plt.title("BiGRU – Training vs Validation Loss")
plt.show()

```



Step 15 — Evaluate BiGRU and Compare with BiLSTM

We compute metrics in:

- scaled space
- original kW units

Then we print a side-by-side comparison table to pick the final best model for Flask.

```
# Load best BiGRU model
best_bigru = keras.models.load_model(
    "/kaggle/working/artifacts/model_cnn_bigru_sa.keras",
    custom_objects={"SelfAttention": SelfAttention}
)

# Predict
y_pred_bigru = best_bigru.predict(X_test, verbose=0)

# Scaled metrics
rmse_bigru = math.sqrt(mean_squared_error(y_test, y_pred_bigru))
mae_bigru = mean_absolute_error(y_test, y_pred_bigru)
r2_bigru = r2_score(y_test, y_pred_bigru)

# In original kW
y_pred_bigru_kw = inverse_target(scaler, y_pred_bigru, num_cols,
t_idx)

rmse_bigru_kw = math.sqrt(mean_squared_error(y_test_kw,
y_pred_bigru_kw))
mae_bigru_kw = mean_absolute_error(y_test_kw, y_pred_bigru_kw)
r2_bigru_kw = r2_score(y_test_kw, y_pred_bigru_kw)

# Comparison
comp = pd.DataFrame({
    "Model": ["CNN-BiLSTM-SA", "CNN-BiGRU-SA"],
    "RMSE_scaled": [rmse, rmse_bigru],
    "MAE_scaled": [mae, mae_bigru],
    "R2_scaled": [r2, r2_bigru],
    "RMSE_kw": [rmse_kw, rmse_bigru_kw],
    "MAE_kw": [mae_kw, mae_bigru_kw],
    "R2_kw": [r2_kw, r2_bigru_kw]
})

display(comp)

# Save BiGRU metrics too
metrics_bigru = {
    "scaled": {"rmse": float(rmse_bigru), "mae": float(mae_bigru),
    "r2": float(r2_bigru)},
    "original_kw": {"rmse": float(rmse_bigru_kw), "mae": float(mae_bigru_kw),
    "r2": float(r2_bigru_kw)}
}
```

```

        float(mae_bigru_kw), "r2": float(r2_bigru_kw)} ,
        "selected_features": selected_features,
        "lookback": LOOKBACK,
        "horizon": HORIZON
    }

with open("/kaggle/working/artifacts/metrics_cnn_bigru_sa.json", "w")
as f:
    json.dump(metrics_bigru, f, indent=2)

print("□ Saved: /kaggle/working/artifacts/metrics_cnn_bigru_sa.json")

    Model   RMSE_scaled   MAE_scaled   R2_scaled   RMSE_kw
MAE_kw \
0  CNN-BiLSTM-SA      0.072253     0.050977     0.562507     0.465060
0.328117
1  CNN-BiGRU-SA       0.073396     0.052618     0.548556     0.472417
0.338677

    R2_kw
0  0.562507
1  0.548556

□ Saved: /kaggle/working/artifacts/metrics_cnn_bigru_sa.json

```

Step 16 — Finalize Best Model Artifacts

We pick the best-performing model: **CNN-BiLSTM-SA**.

We create a clean final artifact set for VS Code integration:

- final_model.keras
- scaler.pkl
- selected_features.json
- config.json
- metrics_final.json

```

import shutil, glob

ART = "/kaggle/working/artifacts"
os.makedirs(ART, exist_ok=True)

# Copy best model as final alias (easier for Flask)
src_model = f"{ART}/model_cnn_bilstm_sa.keras"
final_model = f"{ART}/final_model.keras"
shutil.copyfile(src_model, final_model)

# Copy BiLSTM metrics as final metrics
src_metrics = f"{ART}/metrics_cnn_bilstm_sa.json"
final_metrics = f"{ART}/metrics_final.json"

```

```

shutil.copyfile(src_metrics, final_metrics)

print("□ Final artifacts created:")
print(" -", final_model)
print(" -", final_metrics)

# List what we have
print("\n□ artifacts folder contents:")
for p in sorted(glob.glob(f"{ART}/*")):
    print(" -", os.path.basename(p))

□ Final artifacts created:
- /kaggle/working/artifacts/final_model.keras
- /kaggle/working/artifacts/metrics_final.json

□ artifacts folder contents:
- config.json
- final_model.keras
- metrics_cnn_bigru_sa.json
- metrics_cnn_bilstm_sa.json
- metrics_final.json
- model_cnn_bigru_sa.keras
- model_cnn_bilstm_sa.keras
- scaler.pkl
- selected_features.json

```

Step 17 — Zip Artifacts for Download

Kaggle makes it easiest if everything is packaged into a single ZIP. We generate:

- `electricity_forecasting_artifacts.zip`

Then you download from Kaggle Output panel.

```

import zipfile

zip_path = "/kaggle/working/electricity_forecasting_artifacts.zip"
files_to_zip = [
    "/kaggle/working/artifacts/final_model.keras",
    "/kaggle/working/artifacts/scaler.pkl",
    "/kaggle/working/artifacts/selected_features.json",
    "/kaggle/working/artifacts/config.json",
    "/kaggle/working/artifacts/metrics_final.json",
]

with zipfile.ZipFile(zip_path, "w", compression=zipfile.ZIP_DEFLATED) as z:
    for fpath in files_to_zip:
        if os.path.exists(fpath):
            z.write(fpath, arcname=os.path.basename(fpath))

```

```

    else:
        print("⚠ Missing (not zipped):", fpath)

print("ZIP created:", zip_path)
print("→ Kaggle: Right sidebar → Output → Download")
print("electricity_forecasting_artifacts.zip")

ZIP created: /kaggle/working/electricity_forecasting_artifacts.zip
→ Kaggle: Right sidebar → Output → Download
electricity_forecasting_artifacts.zip

```

Step 18 — Quick Inference Test (Flask Simulation)

This verifies that:

- model loads correctly
- scaler + feature order works
- prediction pipeline runs using the latest `lookback` window

We simulate how Flask will work:

- take last 24 rows
- scale using saved scaler
- predict next-hour `Global_active_power`
- inverse transform to kW

```

# Load final artifacts
scaler_loaded = joblib.load("/kaggle/working/artifacts/scaler.pkl")
with open("/kaggle/working/artifacts/selected_features.json") as f:
    feats = json.load(f)
with open("/kaggle/working/artifacts/config.json") as f:
    cfg = json.load(f)

lookback = int(cfg["lookback"])
target_col = cfg["target_col"]
all_cols_live = feats + [target_col]

# Load final model
final_loaded = keras.models.load_model(
    "/kaggle/working/artifacts/final_model.keras",
    custom_objects={"SelfAttention": SelfAttention}
)

# Prepare last window from df_hourly (unscaled)
last_block = df_hourly[all_cols_live].iloc[-lookback:].values # (24, 6)
last_scaled = scaler_loaded.transform(last_block) # (24, 6)

X_live = last_scaled.reshape(1, lookback, last_scaled.shape[1]) #

```

```
(1,24,6)

pred_scaled = final_loaded.predict(X_live, verbose=0).reshape(-1, 1)

# Inverse-transform to kW using dummy array trick
num_cols_live = len(all_cols_live)
t_idx_live = num_cols_live - 1

pred_kw = inverse_target(scaler_loaded, pred_scaled, num_cols_live,
t_idx_live)[0]

print("□ Flask-style inference test passed.")
print("Predicted next-hour Global_active_power (kW):", float(pred_kw))

□ Flask-style inference test passed.
Predicted next-hour Global_active_power (kW): 0.7821608781814575

# □ SINGLE-CELL: Generate ALL thesis figures + ZIP download (Kaggle-ready)
# Assumes these exist from your previous steps:
# history, history_bigru, y_test, y_pred, y_test_kw, y_pred_kw, mi_df
# If comp doesn't exist, it will be created from your metrics.

import os, json, glob, zipfile
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt

FIG_DIR = "/kaggle/working/figures"
os.makedirs(FIG_DIR, exist_ok=True)

def savefig(name):
    path = os.path.join(FIG_DIR, name)
    plt.tight_layout()
    plt.savefig(path, dpi=300, bbox_inches="tight")
    plt.close()
    print("□ Saved:", path)

# -----
# Build/ensure comp DataFrame
# -----
try:
    comp # if already exists from your notebook
except NameError:
    comp = pd.DataFrame({
        "Model": ["CNN-BiLSTM-SA", "CNN-BiGRU-SA"],
        "RMSE_scaled": [0.072253, 0.073396],
        "MAE_scaled": [0.050977, 0.052618],
        "R2_scaled": [0.562507, 0.548556],
        "RMSE_kW": [0.465060, 0.472417],
```

```

        "MAE_kw":      [0.328117, 0.338677],
        "R2_kw":       [0.562507, 0.548556],
    })

# -----
# 1) Model comparison table image
# -----
fig, ax = plt.subplots(figsize=(12, 2.2))
ax.axis("off")
tbl = ax.table(
    cellText=np.round(comp.drop(columns=["Model"]).values, 6),
    colLabels=comp.drop(columns=["Model"]).columns,
    rowLabels=comp["Model"].values,
    loc="center"
)
tbl.auto_set_font_size(False)
tbl.set_fontsize(9)
tbl.scale(1, 1.4)
plt.title("Model Performance Comparison", pad=10)
savefig("Fig_Model_Comparison_Table.png")

# -----
# 2) Model comparison bar charts
# -----
models = comp["Model"].values

plt.figure(figsize=(7,4))
plt.bar(models, comp["RMSE_kw"].values)
plt.ylabel("RMSE (kW)")
plt.title("RMSE Comparison (kW)")
savefig("Fig_Compare_RMSE_kw.png")

plt.figure(figsize=(7,4))
plt.bar(models, comp["MAE_kw"].values)
plt.ylabel("MAE (kW)")
plt.title("MAE Comparison (kW)")
savefig("Fig_Compare_MAE_kw.png")

plt.figure(figsize=(7,4))
plt.bar(models, comp["R2_kw"].values)
plt.ylabel("R2")
plt.title("R2 Comparison (kW)")
savefig("Fig_Compare_R2_kw.png")

# -----
# 3) Training curves (BiLSTM + BiGRU)
# -----
# BiLSTM
plt.figure(figsize=(7,4))
plt.plot(history.history["loss"], label="train_loss")

```

```

plt.plot(history.history["val_loss"], label="val_loss")
plt.xlabel("Epoch")
plt.ylabel("MSE Loss")
plt.title("CNN-BiLSTM-SA – Training vs Validation Loss")
plt.legend()
savefig("Fig_Loss_CNN_BiLSTM_SA.png")

# BiGRU
plt.figure(figsize=(7,4))
plt.plot(history_bigru.history["loss"], label="train_loss")
plt.plot(history_bigru.history["val_loss"], label="val_loss")
plt.xlabel("Epoch")
plt.ylabel("MSE Loss")
plt.title("CNN-BiGRU-SA – Training vs Validation Loss")
plt.legend()
savefig("Fig_Loss_CNN_BiGRU_SA.png")

# -----
# 4) Actual vs Predicted (Scaled + kW) – first N
# -----
N = 300
y_test_s = np.array(y_test).reshape(-1)
y_pred_s = np.array(y_pred).reshape(-1)
y_test_kw_s = np.array(y_test_kw).reshape(-1)
y_pred_kw_s = np.array(y_pred_kw).reshape(-1)

plt.figure(figsize=(10,4))
plt.plot(y_test_s[:N], label="Actual")
plt.plot(y_pred_s[:N], label="Predicted")
plt.xlabel("Time step")
plt.ylabel("Scaled Global_active_power")
plt.title("Actual vs Predicted (Scaled) – First 300 Test Points")
plt.legend()
savefig("Fig_Actual_vs_Pred_Scaled_First300.png")

plt.figure(figsize=(10,4))
plt.plot(y_test_kw_s[:N], label="Actual (kW)")
plt.plot(y_pred_kw_s[:N], label="Predicted (kW)")
plt.xlabel("Time step")
plt.ylabel("Global_active_power (kW)")
plt.title("Actual vs Predicted (kW) – First 300 Test Points")
plt.legend()
savefig("Fig_Actual_vs_Pred_kW_First300.png")

# -----
# 5) Residual analysis (kW)
# -----
res = y_test_kw_s - y_pred_kw_s

M = min(1000, len(res))

```

```

plt.figure(figsize=(10,4))
plt.plot(res[:M])
plt.axhline(0)
plt.xlabel("Time step")
plt.ylabel("Residual (kW)")
plt.title("Residuals Over Time (kW) – First 1000 Test Points")
savefig("Fig_Residuals_Time_kw.png")

plt.figure(figsize=(7,4))
plt.hist(res, bins=60)
plt.xlabel("Residual (kW)")
plt.ylabel("Frequency")
plt.title("Residual Distribution (kW)")
savefig("Fig_Residuals_Hist_kw.png")

# -----
# 6) Scatter: Predicted vs Actual (kW)
# -----
plt.figure(figsize=(6,6))
plt.scatter(y_test_kw_s, y_pred_kw_s, s=6, alpha=0.4)
mn = float(min(y_test_kw_s.min(), y_pred_kw_s.min()))
mx = float(max(y_test_kw_s.max(), y_pred_kw_s.max()))
plt.plot([mn, mx], [mn, mx], linestyle="--")
plt.xlabel("Actual (kW)")
plt.ylabel("Predicted (kW)")
plt.title("Predicted vs Actual (kW)")
savefig("Fig_Scatter_Pred_vs_Actual_kw.png")

# -----
# 7) Feature relevance plot (MI ranking)
# -----
# If mi_df is missing, we create it from your selected MI scores
# (based on your output).
try:
    mi_df
except NameError:
    mi_df = pd.DataFrame({
        "feature": ["Global_intensity", "Sub_metering_3", "Voltage",
        "Global_reactive_power", "Sub_metering_2", "Sub_metering_1"],
        "MI": [3.242534, 0.655366, 0.202940, 0.173717, 0.172238,
        0.153720]
    })

plt.figure(figsize=(8,4))
plt.bar(mi_df["feature"].values, mi_df["MI"].values)
plt.xticks(rotation=20, ha="right")
plt.ylabel("Mutual Information")
plt.title("Feature Relevance (Mutual Information vs Target)")
savefig("Fig_Feature_MI_Ranking.png")

```

```
# -----
# 8) ZIP all figures
# -----
zip_fig = "/kaggle/working/thesis_figures.zip"
with zipfile.ZipFile(zip_fig, "w", compression=zipfile.ZIP_DEFLATED) as z:
    for p in sorted(glob.glob(f"{FIG_DIR}/*.png")):
        z.write(p, arcname=os.path.basename(p))

print("\n\s All figures generated in:", FIG_DIR)
print("\s Zipped figures:", zip_fig)
print("→ Kaggle Output panel → Download thesis_figures.zip")

\s Saved: /kaggle/working/figures/Fig_Model_Comparison_Table.png
\s Saved: /kaggle/working/figures/Fig_Compare_RMSE_kw.png
\s Saved: /kaggle/working/figures/Fig_Compare_MAE_kw.png
\s Saved: /kaggle/working/figures/Fig_Compare_R2_kw.png
\s Saved: /kaggle/working/figures/Fig_Loss_CNN_BiLSTM_SA.png
\s Saved: /kaggle/working/figures/Fig_Loss_CNN_BiGRU_SA.png
\s Saved:
/kaggle/working/figures/Fig_Actual_vs_Pred_Scaled_First300.png
\s Saved: /kaggle/working/figures/Fig_Actual_vs_Pred_kw_First300.png
\s Saved: /kaggle/working/figures/Fig_Residuals_Time_kw.png
\s Saved: /kaggle/working/figures/Fig_Residuals_Hist_kw.png
\s Saved: /kaggle/working/figures/Fig_Scatter_Pred_vs_Actual_kw.png
\s Saved: /kaggle/working/figures/Fig_Feature_MI_Ranking.png

\s All figures generated in: /kaggle/working/figures
\s Zipped figures: /kaggle/working/thesis_figures.zip
→ Kaggle Output panel → Download thesis_figures.zip
```